





A software does not just provide APIs, it also uses them

- We could not worry about the APIs that it uses.
- We'd like the APIs to do the right thing, but we need to test this
- But don't have direct control, sometimes we can simulate what we want, but it's not always the case

A way to deal with this is using low level programming languages or interfaces that are predictable and return friendlier error code when we have the option (eg between UNIX system calls and python libraries, you almost always use the python libraries)

Another way is to use fault injection

`file = open("/tmp/foo", 'w')`

`file = my_open("/tmp/foo", 'w')`

you may have a hard time finding testing where the open call fails. /tmp is in most machines

A stub function, almost functionally equivalent to open, wrote by us. Conditionally, inside that function's code, we can cause the call to open to fail

A fault injection should not make the code fail to often

### Timing

All those we talked about do not represent all possible inputs and outputs that we might care about. A really common example is timing of inputs.

To know if you need to test timing think of the specification, about what it's trying to do, about its requirements. And also, look at the source code for timeout, sleeps or values that depend on the time.

### Non-Functional inputs

Inputs that affects the operations of a software under test, that have nothing to do with the APIs provided and used by the software we're testing

- context-switches

depending on thread scheduling, bugs can be either concealed or revealed.

The timing is not controlled by our application. This makes testing multithreaded software difficult

### Short survey

White box testing. The tester is using detailed knowledge about the internals of the system in order to construct better test cases.

Black box. We're testing based on how it is supposed to respond

Unit testing. Look at small software modules at a time, and testing it in a isolated fashion. Can be either white box or black box testing.

The goal is to find defects on the internal logic of the S.U.T. as early as possible in order to create more robust software modules.

At this level, we have no hypothesis about the patterns of the S.U.T., we try to test the unit with inputs from all different parts of the domain.

There are many frameworks for unit testing

System testing, we want to know whether the system as a whole meets his goal. Often at this point we're doing black box testing on - the system is now large enough that a knowledge of its internals does not really help creating good test cases

Stress testing, the system is tested at or beyond its normal usage limits. To assess robustness and reliability of the SUT

Random testing, randomly create test inputs. Useful for finding corner cases

### Being great at testing

- developer: "I want this code to succeed"
- tester: "I want this code to fail"

as the tester end goal is creating stronger code which later on doesn't fail

doublethink involved to be both a great tester and developer.

The contradictory nature of those beliefs is apparent as they both have the end goal of creating great software

- Learn to test creatively

(does a much better job than rote testing. Even worse rote testing with obvious inputs)

- Don't ignore weird things. we usually get little hints of things wrong software that would've lead us to find issues



## COVERAGE TESTING

You can think of problems in released software as coming from things that people forgot to test. That means that testing was inadequate but developers were unaware of that.

**Code coverage** is a collection of techniques where automated tools can tell us places where our testing strategy is not doing a good job.

### How much testing is enough?



We want some sort of automated system that looks at our testing effort and gives us a score.

This helps us find parts of the input domain that needs more testing.

Also, we might be able to argue that we've done enough testing, or that the system hasn't been tested enough yet.

### Partitioning the input domain

You want to partition the input domain into a number of different classes, so that all of the points within each class are treated the same by the SUT. While finding those classes we could look at specs, implementation or use our vague intuition.

our testing is confined to some small part of the input domain. Even the small part may contain an infinite # of test cases. There are other parts of the input domain we didn't think to test, that may result in outputs that are not okay. It depends on how you've broken up the input domain.

It would be nice if we could take a large test suite, one that maybe takes several days to run, and identify parts that are completely redundant, that have roughly the same testing effect on the system. Assigning a score lets us do that as well.

Sometimes, what we thought was a class of inputs that are all equivalent, isn't really, and there is a different class hiding within this class which triggers a bug, even though the original test cases didn't.

So in practice what we ended up with is not this idea of coming up with a good partitioning for the input domain. Rather the notion of test coverage.

Test coverage is trying to accomplish the same thing that partitioning was accomplishing, but it does it in a different way. It is an automatic way of partitioning the input domain, based on observed features of the source code.

For example, a particular kind of test coverage is function coverage, which is achieved if every function in our source code is executed during testing.



We end up with a score, called test coverage metrics.

eg 181/250 functions

that would not be a good score, we can look at those functions and come up with a test input that will cause to execute them.

If we cannot find those inputs it either can't be caused or we don't know the system well enough.

Test coverage, a measure of the proportion of the program exercised during testing.

+ gives us an objective score  
+ when coverage is < 100%, we are given meaningful tasks.

- not very helpful in filling errors of omission ... since it's white box.  
things that we should've implemented

- difficult to interpret scores < 100%.  
For large, complex software system, achieving this is often difficult.

- 100% coverage doesn't mean all bugs were found



"When our test coverage fails, it is giving us a bit of evidence, that our test suite is poorly thought out, meaning it is failing to exercise functionality present in our code."

"We should not complete coverage and a number of successful test cases fool us into thinking that a piece of code is right. Often deeper analysis is necessary"

## COVERAGE METRICS

there are a lot of them!

Statement coverage,  
Must execute all statements

Line coverage SIMILAR TO PREVIOUS  
Must execute every line in the source code (a line can have multiple lines)

Branch coverage, a branch in the code is covered if it executes both ways

```
if x == 0:
    y += 1
    if y == 0:
        x += 1
```

← to achieve 100% branch coverage here, you want to execute the if's and the else's

Other coverage metrics are interesting, even though you may not use them in practice, because they form part of the answer to the question "how shall we come up with good test inputs, in order to find bugs?"

loop coverage, executes each loop zero, one more than anytime

loop boundary conditions, are extremely frequent source of bugs in real codes.

Modified condition decision coverage For safety critical software

MCD, branch coverage + conditions  
take all possible values + every condition independently affects its outcome

path coverage, covers how you get to a certain piece of code.

```
def foo(x, y):
    for i in range(x):
        something(i)
    if y:
        somethingElse(i)
```

As x increases, there are unlimited branches. But it gives something to think about when generating test-cases

x, 0 y, true x, 1 y, true x, 2 y, false

Boundary coverage, when a program depends on some numerical range, and when a program has different behaviours based on numbers within that range, then we should test numbers close to the boundary

What about concurrent software?

```
xfer(amount)
shared → a1 += amount
between → a2 -= amount
different → return OK
ways
```

As long as we transfer money between different accounts, probably everything is alright

since the function does not use locks while it manipulates the accounts, if multiple threads operate on the same account concurrently, then it's going to be a problem

Our coverage metrics makes sure that threads T1 and T2, both call the function at the same time while transferring between the same amount. T1 got part way into the function and stops running, and then we start T2 and resume T1 after it finishes